

## Files

- notebooks/analyze\_stacking\_regressor\_vs\_voting\_regressor.py
- local/eval\_results/experiment\_stacking\_regressor\_vs\_voting\_regressor.csv
- experiments/experiment\_stacking\_regressor\_vs\_voting\_regressor.py

## Motivation

The regression method that we currently use for the BucketRegressor is based on voting. A **voting regressor** is an ensemble meta-estimator that fits several base regressors, each on the whole dataset. Then it provides a **weighted** average of the individual predictions to form a final prediction.

While this regression method proved efficient in the past, we wanted to try a new formula: a **stacking regressor**. The primary idea of stacking is to feed the predictions of numerous base models into a higher-level model known as the meta-model or blender, which then combines them to get the final forecast. The architectural advantage of this model is that the stacking process can happen on multiple layers: the blender of the first layer becomes the final model of the second layer etc.

Sklearn provides this implementation of the stacking regressor: <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.StackingRegressor.html>

Here is an example of generalised (multi-layered) stacking: [https://scikit-learn.org/stable/auto\\_examples/ensemble/plot\\_stack\\_predictors.html#sphx-glr-auto-examples-ensemble-plot-stack-predictors-py](https://scikit-learn.org/stable/auto_examples/ensemble/plot_stack_predictors.html#sphx-glr-auto-examples-ensemble-plot-stack-predictors-py)

## Design

Multiple configurations are possible for this stacking type regressor, so we chose two of them which had the greatest potential to provide an accurate comparison with the voting type regressors that were already implemented.

For the implementation of the **simple** stacking regressor we used the following model hierarchy:

Base

- decision tree
- knn
- ridge + kernel approximation (nystroem with rbf)
- sgd + kernel approximation (nystroem with rbf)

Final

- gradient boosting

For the implementation of the **double layered** stacking regressor we used the following model hierarchy:

Base

- decision tree
- knn
- ridge + kernel approximation (nystroem with rbf)
- sgd + kernel approximation (nystroem with rbf)

Finat:

- stacking regressor using:

Base

- extra trees
- svr + kernel approximation (nystroem with rbf)
- random forest

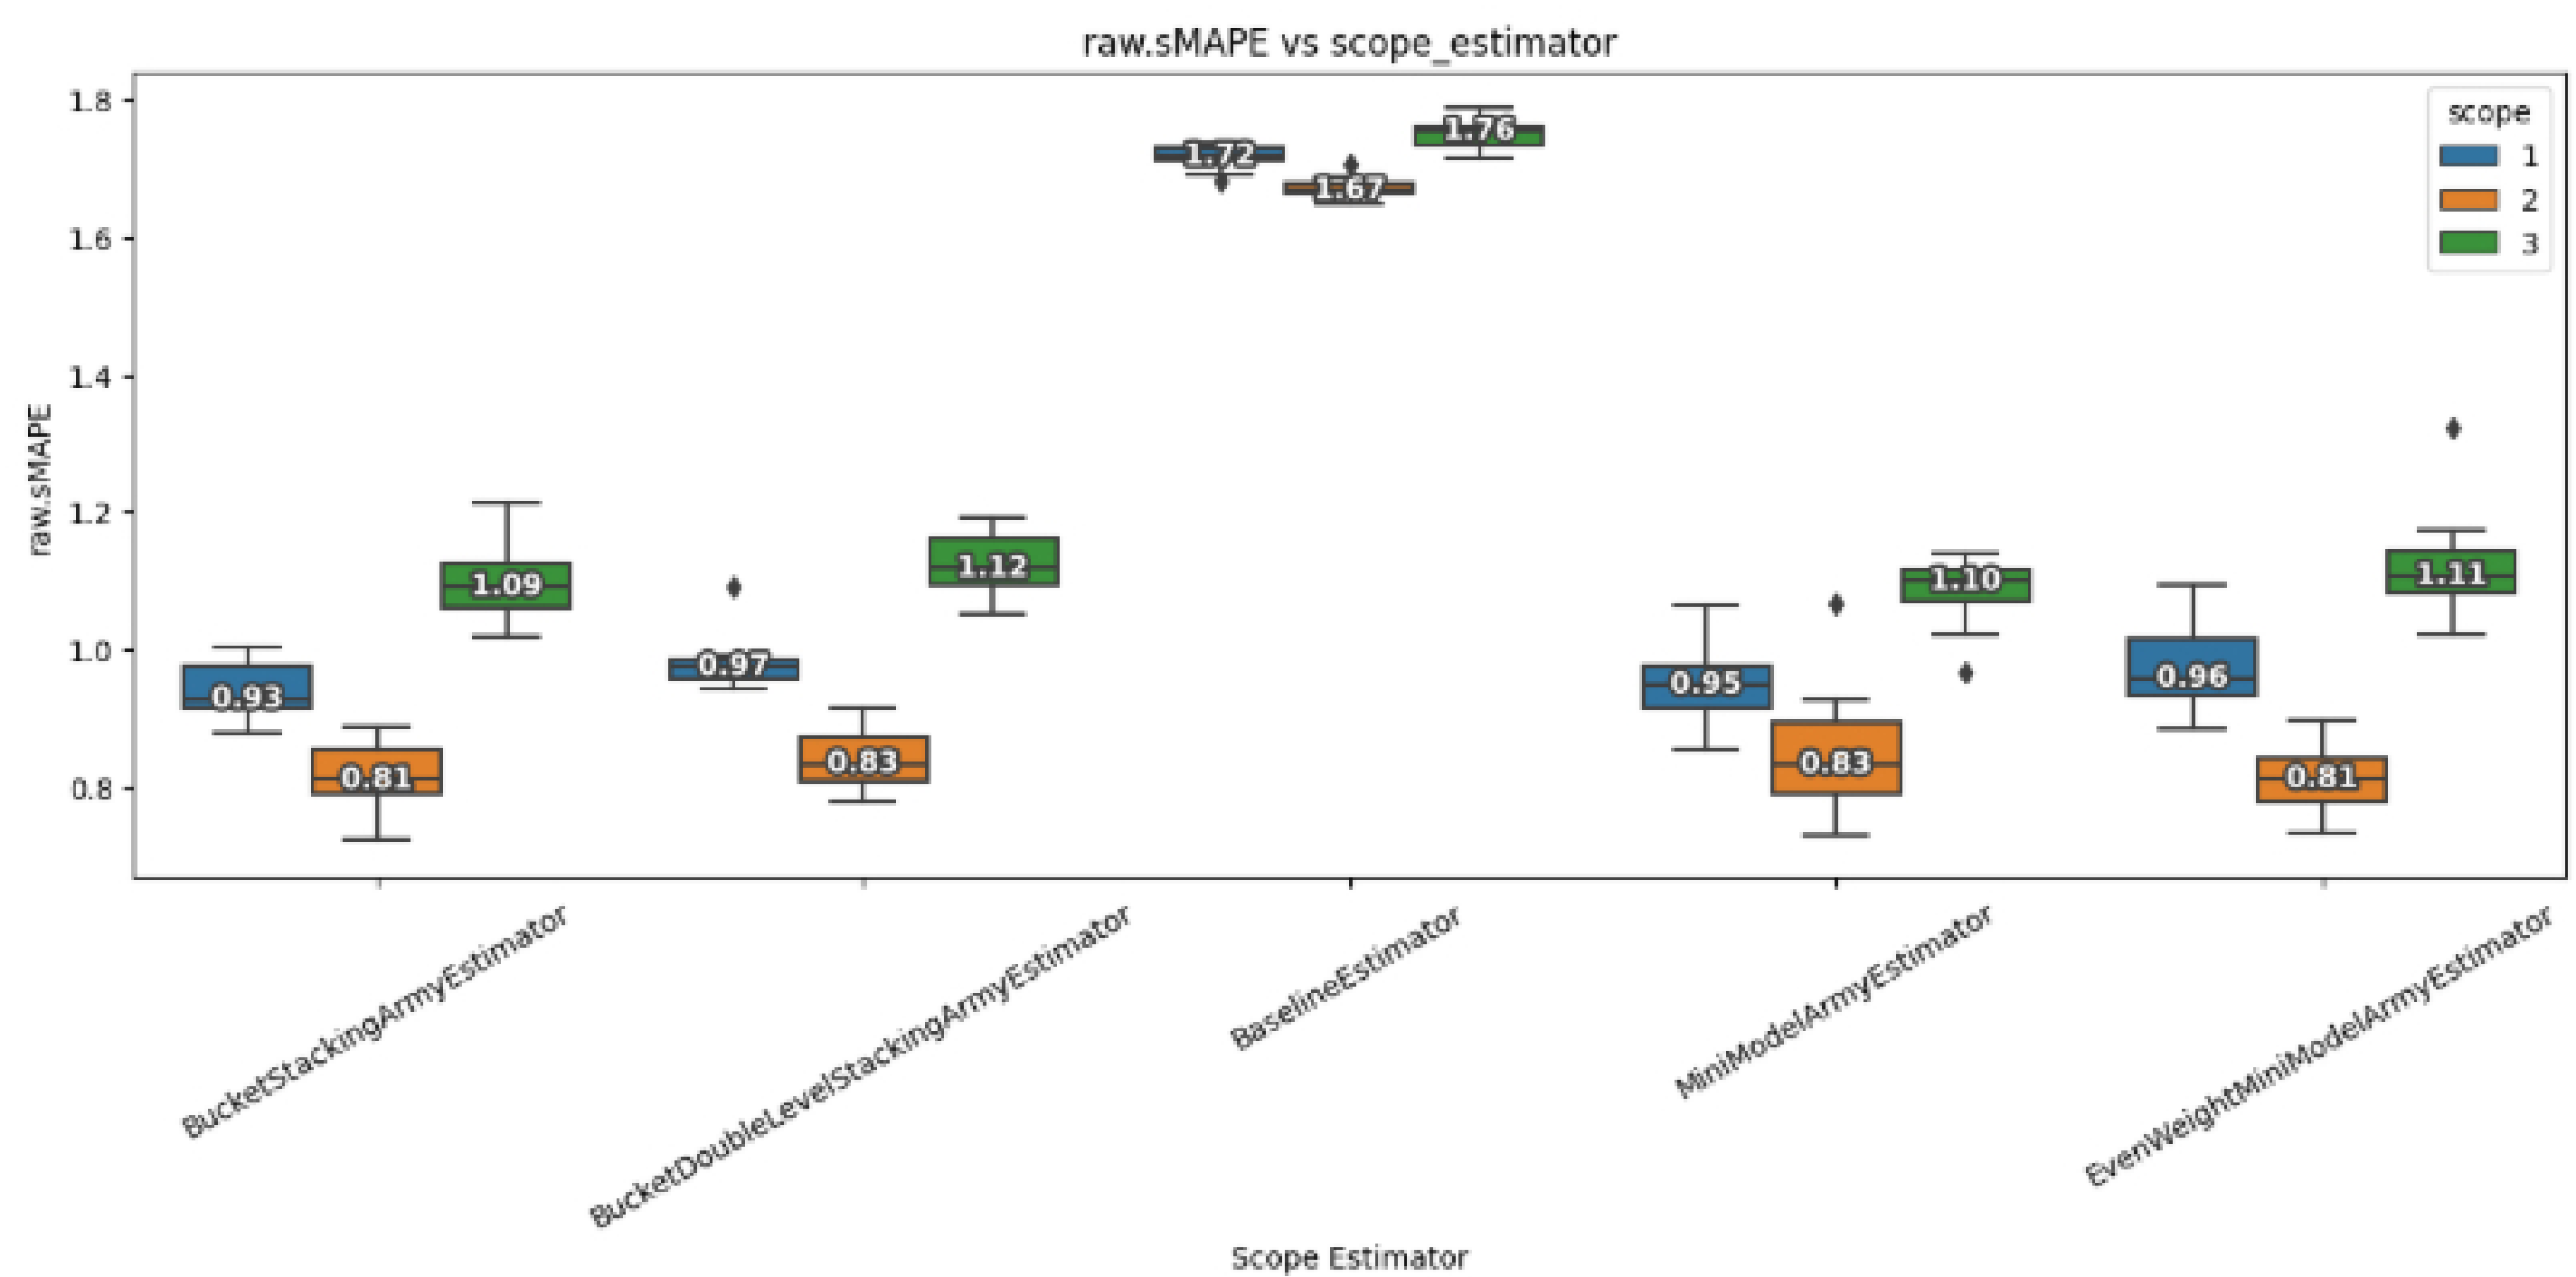
Finat:

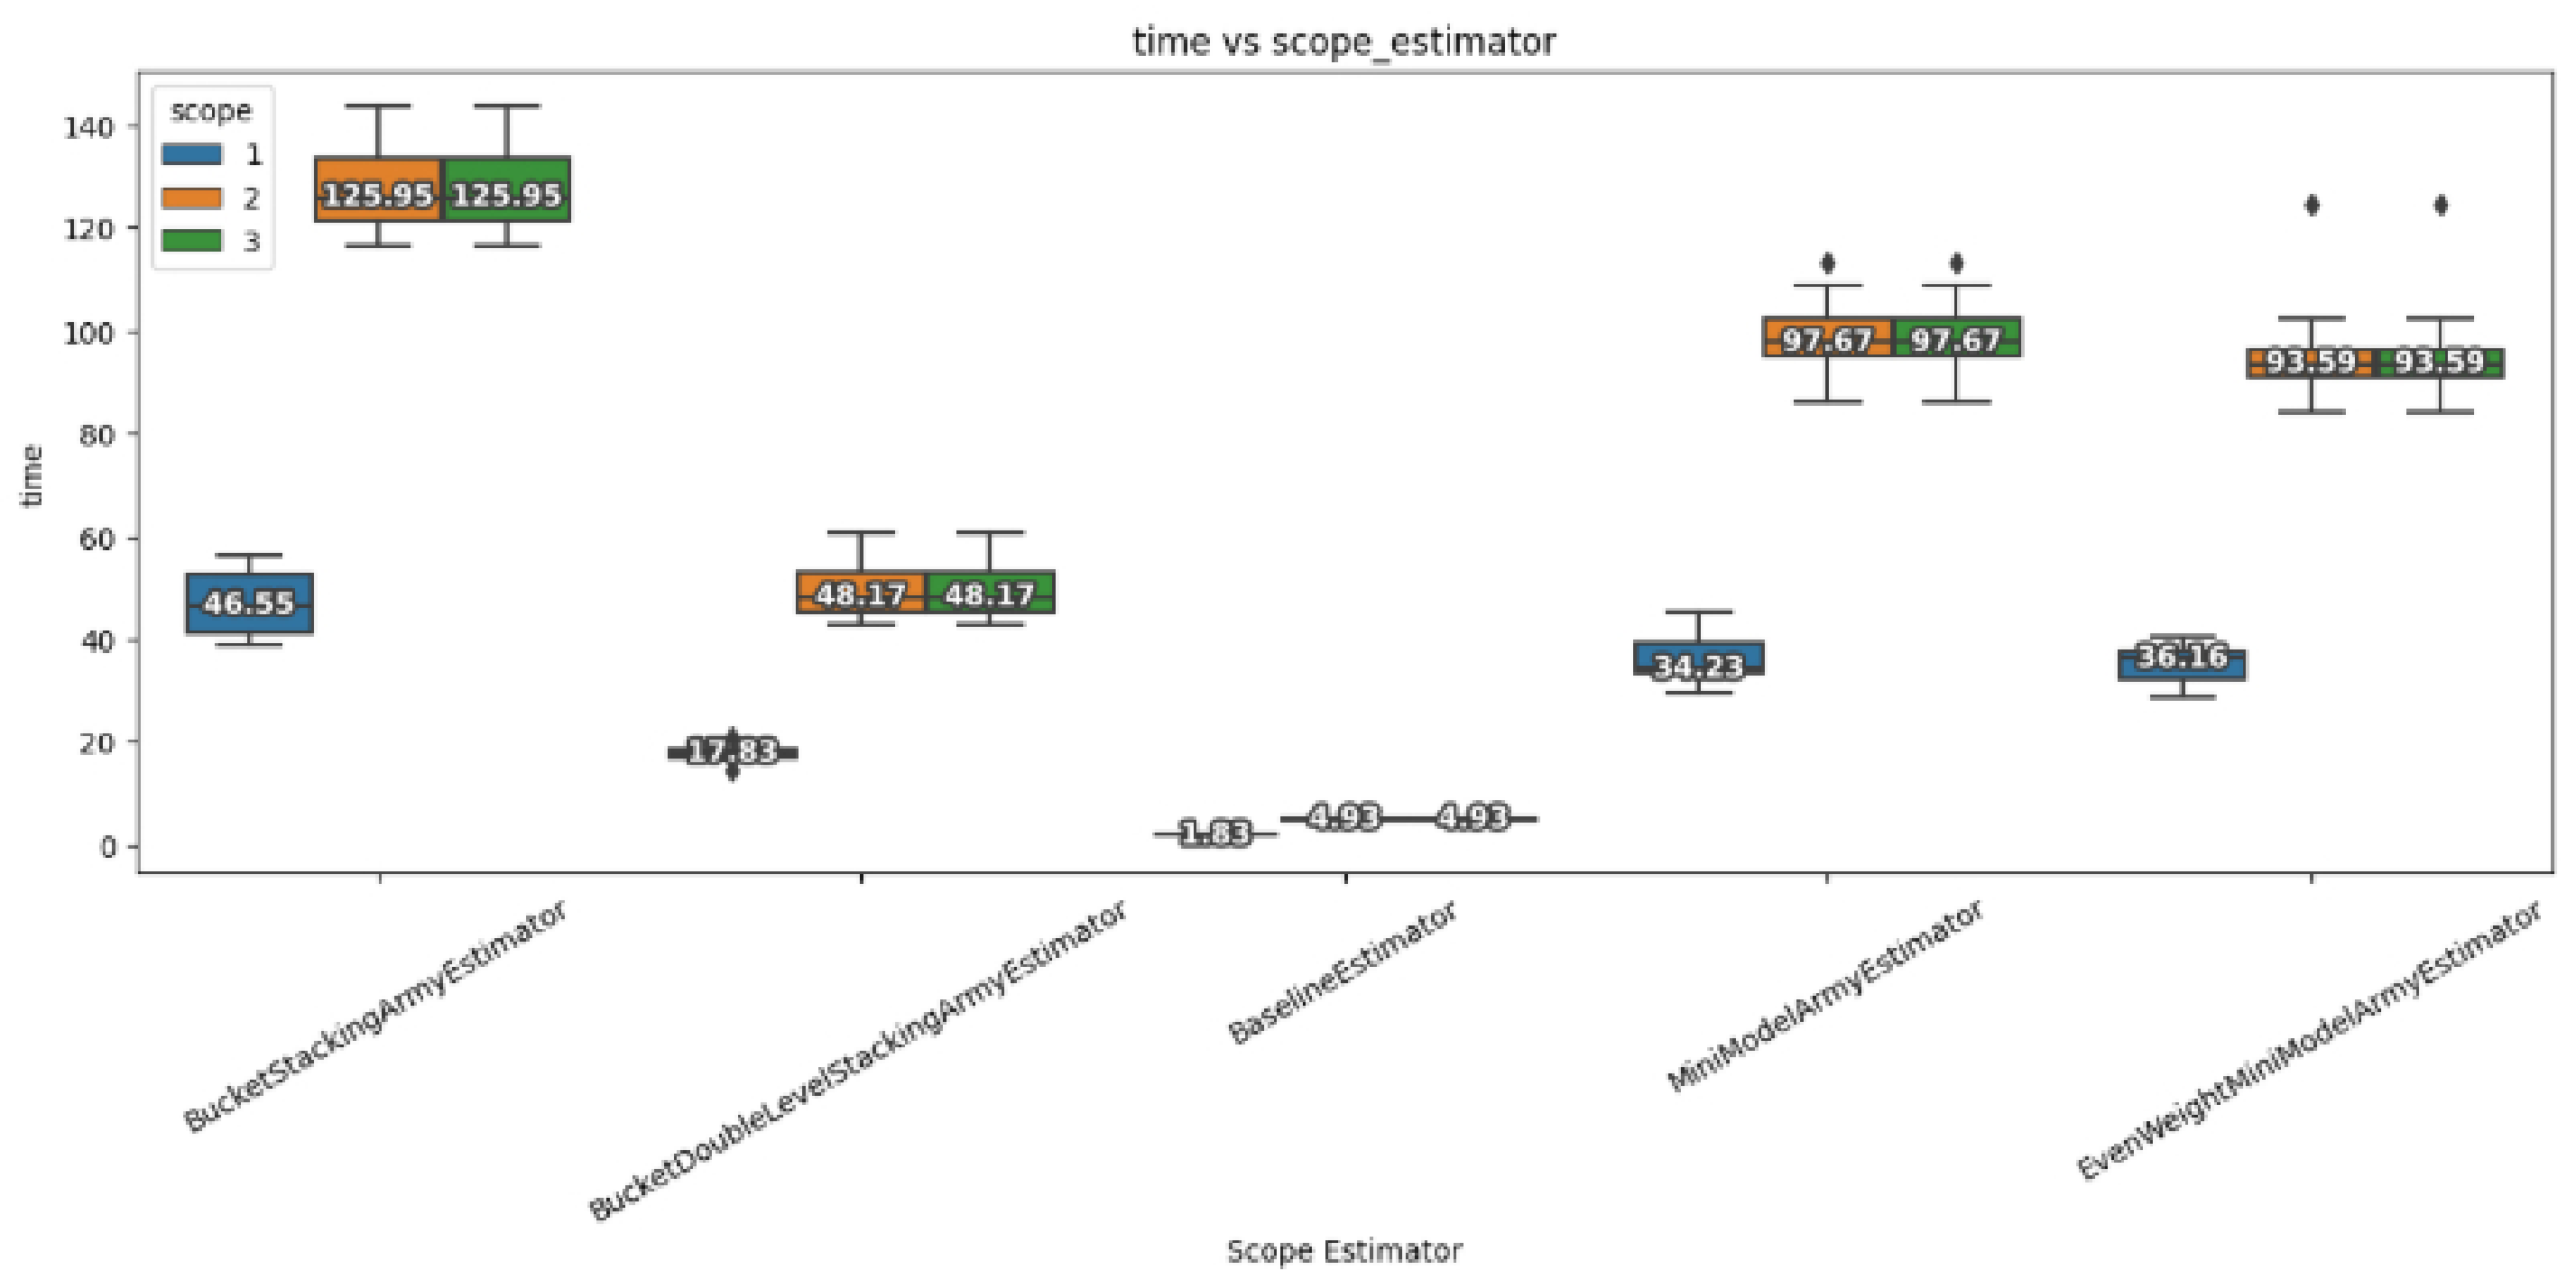
- lgbm

The experiment compares effects on the overall model estimator when using the voting regressor, the even weighted voting regressor, the simple stacking regressor, the double layered stacking regressor, and the baseline one. There are 10 iterations x 5 models, each having a separate data split.

## Results and Insight

Resultsshow that the estimator which usesan even weight voting mechanism for bucket regression (the EvenWeightMiniModelArmyEstimator) still performs best overall. It turns out that the simple stacking strategy improves the sMAPEperformance metric by 3% for scope 1 predictions and 2% for scope 3, but the runtime overhead makes it lessworthy in being considered as a replacement for voting.





We decided to also perform some statistical analysis covering the interaction between scopes and the estimators evaluated above.

The ANOVA table presents a variance analysis between scopes, estimators, and pairs of them. The interaction term between "scope" and "scope\_estimator" explores whether the combination of these two variables has a significant impact. The F-statistic is relatively small, and the p-value (5.699260e-01) is larger than 0.05, indicating that the interaction may not be statistically significant.

ANOVA Table:

	sum_sq	df	F	PR(>F)
C(scope)	1.534977	2.0	206.769715	1.178628e-37
C(scope_estimator)	0.014472	3.0	1.299630	2.783777e-01
C(scope):C(scope_estimator)	0.017877	6.0	0.802704	5.699260e-01
Residual	0.400875	108.0	NaN	NaN

For more details, we decided to look at interactions between specific groups with Tukey's HSD post-hoc test. Unfortunately, the only significant pair results (with  $p < 0.05$ ) are the ones with the baseline as a member of the pair, or the ones that compare the same estimator but for different scopes, which was expected.

Tukey's HSD Post-Hoc Test for Interaction:  
Multiple Comparison of Means - Tukey HSD, FWER=0.05

group1	group2	meandiff	p-adj	lower	upper	reject
1_BaselineEstimator	1_BucketDoubleLevelStackingArmyEstimator	-0.7352	0.0	-0.8204	-0.65	True
1_BaselineEstimator	1_BucketStackingArmyEstimator	-0.7763	0.0	-0.8615	-0.6911	True
1_BaselineEstimator	1_EvenWeightMiniModelArmyEstimator	-0.7462	0.0	-0.8314	-0.661	True
1_BaselineEstimator	1_MiniModelArmyEstimator	-0.7639	0.0	-0.8491	-0.6787	True
1_BaselineEstimator	2_BaselineEstimator	-0.0446	0.8901	-0.1298	0.0406	False
1_BaselineEstimator	2_BucketDoubleLevelStackingArmyEstimator	-0.8749	0.0	-0.9601	-0.7897	True
1_BaselineEstimator	2_BucketStackingArmyEstimator	-0.8997	0.0	-0.9849	-0.8145	True
1_BaselineEstimator	2_EvenWeightMiniModelArmyEstimator	-0.9065	0.0	-0.9917	-0.8213	True
1_BaselineEstimator	2_MiniModelArmyEstimator	-0.8695	0.0	-0.9547	-0.7843	True
1_BaselineEstimator	3_BaselineEstimator	0.0348	0.9851	-0.0504	0.12	False
1_BaselineEstimator	3_BucketDoubleLevelStackingArmyEstimator	-0.5962	0.0	-0.6814	-0.511	True
1_BaselineEstimator	3_BucketStackingArmyEstimator	-0.6201	0.0	-0.7053	-0.5349	True
1_BaselineEstimator	3_EvenWeightMiniModelArmyEstimator	-0.5917	0.0	-0.6769	-0.5065	True
1_BaselineEstimator	3_MiniModelArmyEstimator	-0.6348	0.0	-0.7201	-0.5496	True
1_BucketDoubleLevelStackingArmyEstimator	1_BucketStackingArmyEstimator	-0.0411	0.9395	-0.1263	0.0441	False
1_BucketDoubleLevelStackingArmyEstimator	1_EvenWeightMiniModelArmyEstimator	-0.0109	1.0	-0.0961	0.0743	False
1_BucketDoubleLevelStackingArmyEstimator	1_MiniModelArmyEstimator	-0.0287	0.9978	-0.1139	0.0565	False
1_BucketDoubleLevelStackingArmyEstimator	2_BaselineEstimator	0.6906	0.0	0.6054	0.7758	True
1_BucketDoubleLevelStackingArmyEstimator	2_BucketDoubleLevelStackingArmyEstimator	-0.1397	0.0	-0.2249	-0.0545	True
...						
3_BucketStackingArmyEstimator	3_EvenWeightMiniModelArmyEstimator	0.0284	0.998	-0.0568	0.1136	False
3_BucketStackingArmyEstimator	3_MiniModelArmyEstimator	-0.0147	1.0	-0.1	0.0705	False
3_EvenWeightMiniModelArmyEstimator	3_MiniModelArmyEstimator	-0.0432	0.9129	-0.1284	0.042	False

We also tried the Mixed Linear Model (MixedLM) regression analysis This type of model is used when dealing with clustered or hierarchical data where observations are not independent. In this case, there are 12 groups, and each group has 10 observations. The dependent variable is "raw\_sMAPE" and the independent variable is "scope\_estimator" with multiple levels. The model considers both fixed effects (overall relationships) and random effects (variance within groups).

Fixed effects:

The "coef" column for C(scope\_estimator) shows the coefficients for the different levels of the 'scope\_estimator' variable. Each coefficient represents the estimated change in the dependent variable when the corresponding level of 'scope\_estimator' is compared to the reference level ('BaselineEstimator' in this case). All p-values are below 0.05, suggesting that the variable is not statistically significant. For example:

- 'BucketStackingArmyEstimator': -0.030 (p-value = 0.793)
- 'EvenWeightMiniModelArmyEstimator': -0.013 (p-value = 0.911)
- 'MiniModelArmyEstimator': -0.021 (p-value = 0.856)

Random effects:

Group Var represents the estimated variance of the random effects associated with different groups formed by the interaction of 'scope' and 'scope\_estimator.' In this case, the estimated variance is 0.019 with a standard error of 0.165.

Mixed Linear Model Regression Results						
=====						
Model:	MixedLM	Dependent Variable:		raw_sMAPE		
No. Observations:	120	Method:		REML		
No. Groups:	12	Scale:		0.0037		
Min. group size:	10	Log-Likelihood:		137.3550		
Max. group size:	10	Converged:		Yes		
Mean group size:	10.0					
-----						
		Coef.	Std.Err.	z	P> z	[0.025 0.975]
-----						
Intercept		0.981	0.080	12.197	0.000	0.823 1.139
C(scope_estimator)[T.BucketStackingArmyEstimator]		-0.030	0.114	-0.263	0.793	-0.253 0.193
C(scope_estimator)[T.EvenWeightMiniModelArmyEstimator]		-0.013	0.114	-0.111	0.911	-0.236 0.210
C(scope_estimator)[T.MiniModelArmyEstimator]		-0.021	0.114	-0.181	0.856	-0.244 0.202
Group Var		0.019	0.165			
=====						

## Decision

We will continue using the even weighted voting mechanism because it produces the best performance for the least runtime so far. Further cascading investigations on how much interaction we have between scopes and different estimator configurations might be a good idea in order to improve the overall model performance.